

Class08: Breast Cancer Mini Project

Kyle Canturia (A17502778)

Table of contents

Background	1
Principal Component Analysis	2
Interpreting PCA results	4
Communicating PCA results	6
Hierarchical Clustering	6
Combining Methods	8
Prediction	9

Background

In today's class we will apply the methods and techniques clustering and PCA to help make sense of a real world breast cancer FNA biopsy data set.

```
fna.data <- "WisconsinCancer.csv"  
wisc.df <- read.csv(fna.data, row.names=1)
```

```
wisc.data <- wisc.df[,-1]  
#removes first column from dataset, we don't want to use this for machine learning models, in
```

```
diagnosis <- wisc.df$diagnosis  
#stores diagnosis column from original dataset in variable "diagnosis"
```

Q1: How many observations are in this dataset?

```
View(wisc.data)
```

There are 569 observations in this dataset.

Q2: How many of the observations have a malignant diagnosis?

```
table(wisc.df$diagnosis)
```

```
B      M  
357  212
```

```
sum(wisc.df$diagnosis == "M")
```

```
[1] 212
```

```
table(wisc.df$diagnosis == "M")
```

```
FALSE  TRUE  
357   212
```

There are 212 observations with a malignant diagnosis.

Q3: How many variables/features in the data are suffixed with _mean?

```
sum(grep("_mean$", colnames(wisc.data)))
```

```
[1] 10
```

```
#either works  
length(grep("_mean", colnames(wisc.data)))
```

```
[1] 10
```

There are 10 observations that are suffixed with “mean”.

Principal Component Analysis

The main function here is `prcomp()` and we want to make sure we set the optional argument `scale=TRUE`:

```
wisc.pr <- prcomp(wisc.data, scale = T)
summary(wisc.pr)
```

Importance of components:

	PC1	PC2	PC3	PC4	PC5	PC6	PC7
Standard deviation	3.6444	2.3857	1.67867	1.40735	1.28403	1.09880	0.82172
Proportion of Variance	0.4427	0.1897	0.09393	0.06602	0.05496	0.04025	0.02251
Cumulative Proportion	0.4427	0.6324	0.72636	0.79239	0.84734	0.88759	0.91010
	PC8	PC9	PC10	PC11	PC12	PC13	PC14
Standard deviation	0.69037	0.6457	0.59219	0.5421	0.51104	0.49128	0.39624
Proportion of Variance	0.01589	0.0139	0.01169	0.0098	0.00871	0.00805	0.00523
Cumulative Proportion	0.92598	0.9399	0.95157	0.9614	0.97007	0.97812	0.98335
	PC15	PC16	PC17	PC18	PC19	PC20	PC21
Standard deviation	0.30681	0.28260	0.24372	0.22939	0.22244	0.17652	0.1731
Proportion of Variance	0.00314	0.00266	0.00198	0.00175	0.00165	0.00104	0.0010
Cumulative Proportion	0.98649	0.98915	0.99113	0.99288	0.99453	0.99557	0.9966
	PC22	PC23	PC24	PC25	PC26	PC27	PC28
Standard deviation	0.16565	0.15602	0.1344	0.12442	0.09043	0.08307	0.03987
Proportion of Variance	0.00091	0.00081	0.0006	0.00052	0.00027	0.00023	0.00005
Cumulative Proportion	0.99749	0.99830	0.9989	0.99942	0.99969	0.99992	0.99997
	PC29	PC30					
Standard deviation	0.02736	0.01153					
Proportion of Variance	0.00002	0.00000					
Cumulative Proportion	1.00000	1.00000					

Q4: What proportion of orginal variance is captured by the first principal component (PC1)?

44.27% of the original variance is captured by PC1.

Q5: How many principal components (PCs) are required to describe at least 70% of the original variance in the data?

It takes 3 PCs to describe at least 70% of the original variance.


Q6: How many principal components are required to describe at least 90% of the original variance in the data?

It takes 7 PCs to describe at least 90% of the original variance.

Interpreting PCA results

Q7: What stands out about this plot? Is it easy or difficult to understand and why?

```
biplot(wisc.pr)
```




The main thing that stands out about this plot is that all the clusters are all really close together. This density makes the plot difficult to understand as it's hardd to pick out individual clusters.

Our main PCA “score plot” or “PC plot” of results:


```
library(ggplot2)
```

```
ggplot(wisc.pr$x) +  
  aes(PC1,PC2, col=diagnosis) +  
  geom_point()
```



Q8: Generate a similar plot for PCs 1 and 3, what do you notice about these plots?

```
ggplot(wisc.pr$x) +  
  aes(PC1, PC3, col=diagnosis) +  
  geom_point()
```



I noticed that in this plot, PC1 appears to be flipped relative to the x-plane and is now upside down compared to the first plot. I notice that in both plots, there's still a pretty clear distinction between diagnosis "B" and "M".

Communicating PCA results

Q9: For the first principal component, what is the main component of the loading vector (i.e. `wisc.pr$rotation[,1]`) for the feature `concave.points_mean`? This tells us how much this original feature contributes to the first PC. Are there any features with larger contributions than this one?

Hierarchical Clustering


Q10: What is the height at which the clustering model has 4 clusters?

```
data.scaled <- scale(wisc.data)

data.dist <- dist(data.scaled)

wisc.hclust <- hclust(data.dist, method = "complete")
```

```
plot(wisc.hclust)
abline(h = 19, col="red", lty=2)
```



```
data.dist
hclust (*, "complete")
```

The clustering model has 4 clusters at a height of 19.

```
wisc.hclust.clusters <- cutree(wisc.hclust,k=4)
```

```
table(wisc.hclust.clusters, diagnosis)
```

diagnosis		
wisc.hclust.clusters	B	M
1	12	165
2	2	5
3	343	40
4	0	2

```
table(wisc.hclust.clusters)
```

wisc.hclust.clusters	1	2	3	4
177	7	383	2	


Q12: Which method gives your favorite results for the same data.dist dataset?
Explain your reasoning.

I believe that the `ward.D2` method gives the best results for the same dataset. The other methods connect clusters based on distance, while `ward.D2` groups based on minimizing the variance within a cluster. This helps to dampen some of the noise, which is especially helpful with such a large dataset.

Combining Methods

Here we will take our PCA results and use those as input for clustering. In other words our `wisc.pr$x` scores that we plotted above (the main output from PCA - how the data lie on our new principal component axis/variables) and use a subset of the PCs that capture the most variance as input for `hclust()`.

```
pc.dist <- dist(wisc.pr$x[,1:3])
wisc.pr.hclust <- hclust(pc.dist, method = "ward.D2")
plot(wisc.pr.hclust)
```



pc.dist
hclust (*, "ward.D2")

Cut the dendrogram/tree into two main groups/clusters:

```
grps <- cutree(wisc.pr.hclust, k=2)
table(grps)
```

```
grps
 1   2
203 366
```

```
table(grps, diagnosis)
```

```
diagnosis
grps   B   M
 1   24 179
 2 333  33
```

I want to know how clustering in `grps` with values of 1 or 2 correspond to the expert `diagnosis`

```
table(grps, diagnosis)
```

```
diagnosis
grps   B   M
 1   24 179
 2 333  33
```

My clustering **groups 1** are mostly “M” diagnosis (179) and my clustering **group 2** are mostly “B” diagnosis (333)

24 FP 179 TP 333 TN 33 FN

Prediction

```
#url <- "new_samples.csv"
url <- "https://tinyurl.com/new-samples-CSV"
new <- read.csv(url)
npc <- predict(wisc.pr, newdata=new)
npc
```

```


          PC1        PC2        PC3        PC4        PC5        PC6        PC7
[1,]  2.576616 -3.135913  1.3990492 -0.7631950  2.781648 -0.8150185 -0.3959098
[2,] -4.754928 -3.009033 -0.1660946 -0.6052952 -1.140698 -1.2189945  0.8193031
          PC8        PC9        PC10       PC11       PC12       PC13       PC14
[1,] -0.2307350  0.1029569 -0.9272861  0.3411457  0.375921  0.1610764  1.187882
[2,] -0.3307423  0.5281896 -0.4855301  0.7173233 -1.185917  0.5893856  0.303029
          PC15       PC16       PC17       PC18       PC19       PC20
[1,]  0.3216974 -0.1743616 -0.07875393 -0.11207028 -0.08802955 -0.2495216
[2,]  0.1299153  0.1448061 -0.40509706  0.06565549  0.25591230 -0.4289500
          PC21       PC22       PC23       PC24       PC25       PC26
[1,]  0.1228233  0.09358453  0.08347651  0.1223396  0.02124121  0.078884581
[2,] -0.1224776  0.01732146  0.06316631 -0.2338618 -0.20755948 -0.009833238
          PC27       PC28       PC29       PC30
[1,]  0.220199544 -0.02946023 -0.015620933  0.005269029
[2,] -0.001134152  0.09638361  0.002795349 -0.019015820

```

```

plot(wisc.pr$x[,1:2], col=grps)
points(npc[,1], npc[,2], col="blue", pch=16, cex=3)
text(npc[,1], npc[,2], c(1,2), col="white")

```



```
wisc.pr.hclust.clusters <- cutree(wisc.pr.hclust, k=2)
```

```
table(wisc.pr.hclust.clusters, diagnosis)
```

		diagnosis	
wisc.pr.hclust.clusters	B	M	
1	24	179	
2	333	33	

Q13: How well does the newly created `hclust` model with two clusters separate out the two “M” and “B” diagnoses?

It does well, there is a clear majority of either diagnosis between each cluster with only a small minority of the other diagnosis.

Q14: How well do the hierarchical clustering models you created in the previous sections (i.e. without first doing PCA) do in terms of separating the diagnoses?

```
table(wisc.hclust.clusters, diagnosis)
```

		diagnosis	
wisc.hclust.clusters	B	M	
1	12	165	
2	2	5	
3	343	40	
4	0	2	

They do not do as well for separating the diagnoses, there are multiple clusters instead of 2 clear clusters which have a single diagnosis dominating them.

Q16: Which patient should be prioritized for follow up based on the results?

Patient 2 should be prioritized for follow up based on the results as they reflect a malignant diagnosis.